


CPSC 304 – March 1, 2018

Administrative notes

- Reminder: tutorial this week, but none is due
- Reminder: Midterm #2: March 8@7pm in CIRS 1250 (same room as last time)
 - Closed notes, closed book, closed neighbour
 - Bring ID
 - Will cover through the end of Datalog (I'll put in a stopping point slide)

Let's talk about the project for a minute

- As posted on Piazza, next check point moved to Monday, March 5
- Reminder: these checkpoints are ungrade 
- PHP vs. JDBC
- Group work

Taking it to the next level



Say you're planning a beach vacation

And you wanted to find if it's possible to get
from YVR to OGG (that's on Maui)

Your available information:

Flight(airline,num,origin,destination)


Now what?



A more general Example: Transitive Closure



Suppose we represent a graph w/ relation $Edge(X, Y)$:
 $Edge(a, b)$, $Edge(a, c)$, $Edge(b, d)$, $Edge(c, d)$, $Edge(d, e)$



How can I express the query: *Find all paths*

$Path(X, Y) \text{ :- } Edge(X, Y).$

$Path(X, Y) \text{ :- } Path(X, Z), Path(Z, Y).$

Evaluating Recursive Queries

$Path(X, Y) \text{ :- } Edge(X, Y).$

$Path(X, Y) \text{ :- } Path(X, Z), Path(Z, Y).$

Semantics: evaluate the rules until a *fixed point*:

Iteration #0: Edge: {(a,b), (a,c), (b,d), (c,d), (d,e)}

Path: {}

Iteration #1: Path: {(a,b), (a,c), (b,d), (c,d), (d,e)}

Iteration #2: Path gets the new tuples: (a,d), (b,e), (c,e)

Path: {(a,b), (a,c), (b,d), (c,d), (d,e), (a, d), (b,e), (c, e)}



Iteration #3: Path gets the new tuple: (a,e)

Path: {(a,b), (a,c), (b,d), (c,d), (d,e), (a, d), (b,e), (c, e), (a,e)}

Iteration #4: Nothing changes → Stop.

Note: # of iterations depends on the data. Cannot be anticipated by only looking at the query!

A fun Example



More examples

- Given:
Movie(id, title)
Actor(id, name)
Role(movie-id, actor-id, character)
- Find names of actors who have “Bacon numbers” (assume there’s only one “Kevin Bacon”)

↘
~~CoStars(Aid,Bid):-Role(Mid,Aid,_), Role(Mid,Bid,_)~~
~~CoStars(Aid,Bid):- CoStars(Aid,Cid), CoStars(Cid,Bid)~~
~~Bacon_N(B):-Actor(Aid, "Kevin Bacon"), CoStars(Aid,Bid),~~
~~Actor(Bid,B)~~

Skip the stuff on Magic Sets

- That's Datalog
- It's simple
- It's based on logic
- It's easy to see the join patterns
(especially with anonymous variables)

Learning Goals Revisited



- Given a set of tuples (an input relation) and rules, compute the output relation for a Datalog program.
- Write Datalog programs to query an input relation.
- Explain why we want to extend query languages with recursive queries. Provide good examples of such queries.
- Explain the importance of safe queries, and what makes a Datalog query safe.

Midterm #2 covers through here.

CPSC 304

Introduction to Database Systems

Structured Query Language (SQL)

Textbook Reference
Database Management Systems: Chapter 5

Databases: the continuing saga



When last we left databases...

- We had decided they were great things
- We knew how to conceptually model them in ER diagrams
- We knew how to logically model them in the relational model
- We knew how to normalize our database relations
- We could formally specify queries in Relational Algebra and Datalog

Now: how do most people write queries? SQL!

Learning Goals



- Given the schemas of a relation, create SQL queries using: SELECT, FROM, WHERE, EXISTS, NOT EXISTS, UNIQUE, NOT UNIQUE, ANY, ALL, DISTINCT, GROUP BY and HAVING.
- Show that there are alternative ways of coding SQL queries to yield the same result. Determine whether or not two SQL queries are equivalent.
- Given a SQL query and table schemas and instances, compute the query result.
- Translate a query between SQL and RA.
- Comment on the relative expressive power of SQL and RA.
- Explain the purpose of NULL values and justify their use. Also describe the difficulties added by having nulls.
- Create and modify table schemas and views in SQL.
- Explain the role and advantages of embedding SQL in application programs.
- Write SQL for a small-to-medium sized programming application that requires database access.
- Identify the pros and cons of using general table constraints (e.g., CONSTRAINT, CHECK) and triggers in databases.

Coming up in SQL...

- Data Definition Language (reminder)
- Basic Structure
- Set Operations
- Aggregate Functions
- Null Values
- Nested Subqueries
- Modification of the Database
- Views
- Integrity Constraints
- Putting SQL to work in an application

The SQL Query Language

- Need for a standard since relational queries are used by many vendors
- Consists of several parts:
 - Data Definition Language (DDL)
(a blast from the past (Chapter 3))
 - Data Manipulation Language (DML)
 - Data Query
 - Data Modification

Creating Tables in SQL(DDL) Revisited

- A SQL relation is defined using the **create table** command:

create table r ($A_1 D_1, A_2 D_2, \dots, A_n D_n,$
 (integrity-constraint₁),
 ...,
 (integrity-constraint_k))

- *Integrity constraints can be:*

- *primary and candidate keys*
- *foreign keys*

- Example:

```
CREATE TABLE Student
(sid    CHAR(20),
 name  CHAR(20),
 address CHAR(20),
 phone CHAR(8),
 major  CHAR(4),
PRIMARY KEY (sid))
```


Domain Types in SQL

Reference Sheet

- **char(*n*)**. Fixed length character string with length *n*.
- **varchar(*n*)**. Variable length character strings, with maximum length *n*.
- **int**. Integer (machine-dependent).
- **smallint**. Small integer (machine-dependent).
- **numeric(*p*,*d*)**. Fixed point number, with user-specified precision of *p* digits, with *d* digits to the right of decimal point.
- **real, double precision**. Floating point and double-precision floating point numbers, with machine-dependent precision.
- **float(*n*)**. Floating point number, with user-specified precision of at least *n* digits.
- Null values are allowed in all the domain types.
To prohibit null values declare attribute to be **not null**
- **create domain** in SQL-92 and 99 creates user-defined domain types
create domain *person-name* char(20) not null

Date/Time Types in SQL

Reference Sheet

- **date.** Dates, containing a (4 digit) year, month and date
 - E.g. **date** '2001-7-27'
- **time.** Time of day, in hours, minutes and seconds.
 - E.g. **time** '09:00:30' **time** '09:00:30.75'
- **timestamp:** date plus time of day
 - E.g. **timestamp** '2001-7-27 09:00:30.75'
- **Interval:** period of time
 - E.g. Interval '1' day
 - Subtracting a date/time/timestamp value from another gives an interval value
 - Interval values can be added to date/time/timestamp values
- Relational DBMS offer a variety of functions to
 - extract values of individual fields from date/time/timestamp
 - convert strings to dates and vice versa
 - For instance in Oracle (date is a timestamp):
 - TO_CHAR(date, format)
 - TO_DATE(string, format)
 - format looks like: 'DD-Mon-YY HH:MI.SS'

Running Example (should look familiar)

Movie(MovieID, Title, Year)

StarsIn(MovieID, StarID, Character)

MovieStar(StarID, Name, Gender)

Basic SQL Query

- SQL is based on set and relational operations
- A typical SQL query has the form:

select A_1, A_2, \dots, A_n
from r_1, r_2, \dots, r_m
where P

SELECT	<i>target-list</i>
FROM	<i>relation-list</i>
WHERE	<i>qualification</i>

- A_i s represent attributes
 - r_i s represent relations
 - P is a predicate.
- The result of a SQL query is a table (relation)
- By default, duplicates are not eliminated in SQL relations, which are **bags** or **multisets** and not sets
- Let's compare to relational algebra...

$\pi \rightarrow$ SELECT clause

$\sigma \rightarrow$ WHERE clause

$\bowtie \rightarrow$ FROM and WHERE clause

Basic SQL/RA Comparison example 1

- Find the titles of movies

Basic SQL/RA Comparison example 1

- Find the titles of movies

$\pi_{\text{Title}}(\text{Movie})$

- In SQL, π is in the SELECT clause
- Select only a subset of the attributes

```
SELECT Title  
FROM   Movie
```

- Note duplication can happen!

This is going to be more interesting if you can write your own queries

- Your Oracle (DBMS) accounts on the department servers are ready
- Your login is "ora_[cs ugrad userid]" Your password is "a[your student id]". To login:
 - Log onto remote.ugrad.cs.ubc.ca.
 - Start sqlplus (the interface to Oracle by typing "sqlplus [Oracle login]@ug" ("ug" is the specific database to access)
 - Additional information can be found on Piazza. Note: if you run into problems getting things to work, make sure that:
 - You're logging onto remote.ugrad.cs.ubc.ca
 - Login to the undergraduate machine using your regular user name and password
 - You have the "@ug" at the end of logging in (that says what database to access)
 - you don't forget to prepend your student ID with an "a" for your sql login
- Scripts for most data is on the exercises page

Clicker Question: SQL projection

- Given the table scores:

what is result of
`SELECT Score1,
 Score2
FROM Scores`

Team1	Team2	Score1	Score2
Dragons	Tigers	5	3
Carp	Swallows	4	6
Bay Stars	Giants	2	1
Marines	Hawks	5	3
Ham Fighters	Buffaloes	1	6
Lions	Golden Eagles	8	12

- Which of the following rows is in the answer?

- A. (1,2)
- B. (5,3)
- C. (8,6)
- D. All are in the answer
- E. None are in the answer

clickerprojection.sql

Clicker Question: SQL projection

- Given the table scores:

what is result of
SELECT Score1,
Score2
FROM Scores

- Which of the following rows is in the answer?

Team1	Team2	Score1	Score2
Dragons	Tigers	5	3
Carp	Swallows	4	6
Bay Stars	Giants	2	1
Marines	Hawks	5	3
Ham Fighters	Buffaloes	1	6
Lions	Golden Eagles	8	12

- A. (1,2)
- B. (5,3) **Correct**
- C. (8,6)
- D. All are in the answer
- E. None are in the answer

In SQL, σ is in *Where* clause

```
SELECT *  
FROM Movie  
WHERE Year > 1939
```

You can use:

- attribute names of the relation(s) used in the FROM.

- comparison operators: =, <>, <, >, <=, >=

- apply arithmetic operations: rating*2

- operations on strings (e.g., "||" for concatenation).

- Lexicographic order on strings.

- Pattern matching: s LIKE p

- Special stuff for comparing dates and times.

Basic SQL/RA Comparison example 2

Find female movie stars

Basic SQL/RA Comparison example 2

Find female movie stars

$\sigma_{\text{Gender} = \text{'female'}} \text{MovieStar}$

```
SELECT  *  
FROM    MovieStar  
WHERE   Gender = 'female'
```

Clicker Question: Selection

- Consider Scores(Team, Opponent, RunsFor, RunsAgainst) and query

```
SELECT *  
FROM Scores  
WHERE  
    RunsFor > 5
```

- Which tuple is in the result?

- A. (Swallows, Carp, 6, 4)
- B. (Swallows, Carp, 4)
- C. (12)
- D. (*)

Team	Opponent	RunsFor	RunsAgainst
Dragons	Tigers	5	3
Carp	Swallows	4	6
Bay Stars	Giants	2	1
Marines	Hawks	5	3
Ham Fighters	Buffaloes	1	6
Lions	Golden Eagles	8	12
Tigers	Dragons	3	5
Swallows	Carp	6	4
Giants	Bay Stars	1	2
Hawks	Marines	3	5
Buffaloes	Ham Fighters	6	1
Golden Eagles	Lions	12	8

clickerselection.sql

Clicker Question: Selection

- Consider Scores(Team, Opponent, RunsFor, RunsAgainst) and query

```
SELECT *  
FROM Scores  
WHERE  
    RunsFor > 5
```

- Which tuple is in the result?

- A. (Swallows, Carp, 6, 4)
- B. (Swallows, Carp, 4)
- C. (12)
- D. (*)

Team	Opponent	RunsFor	RunsAgainst
Dragons	Tigers	5	3
Carp	Swallows	4	6
Bay Stars	Giants	2	1
Marines	Hawks	5	3
Ham Fighters	Buffaloes	1	6
Lions	Golden Eagles	8	12
Tigers	Dragons	3	5
Swallows	Carp	6	4
Giants	Bay Stars	1	2
Hawks	Marines	3	5
Buffaloes	Ham Fighters	6	1
Golden Eagles	Lions	12	8

answer A

Selection & Projection – together forever in SQL



We can put these together:

- What are the names of female movie stars?
- What are the titles of movies from prior to 1939?

Selection & Projection – together forever in SQL



We can put these together:

- What are the names of female movie stars?

```
SELECT name  
FROM MovieStar  
WHERE Gender = 'female'
```

- What are the titles of movies from prior to 1939?

```
SELECT title  
FROM Movie  
WHERE year < 1939
```


Selection example (dates)

reserves

SID	BID	Day
22	101	1998-10-10
22	102	1998-10-10
22	103	1998-10-08
22	104	1998-07-10
31	102	1998-11-10
31	103	1998-11-06
31	104	1998-11-12
58	102	1998-11-08
58	103	1998-11-12

Select *
From reserves
Where day < DATE'1998-11-01'

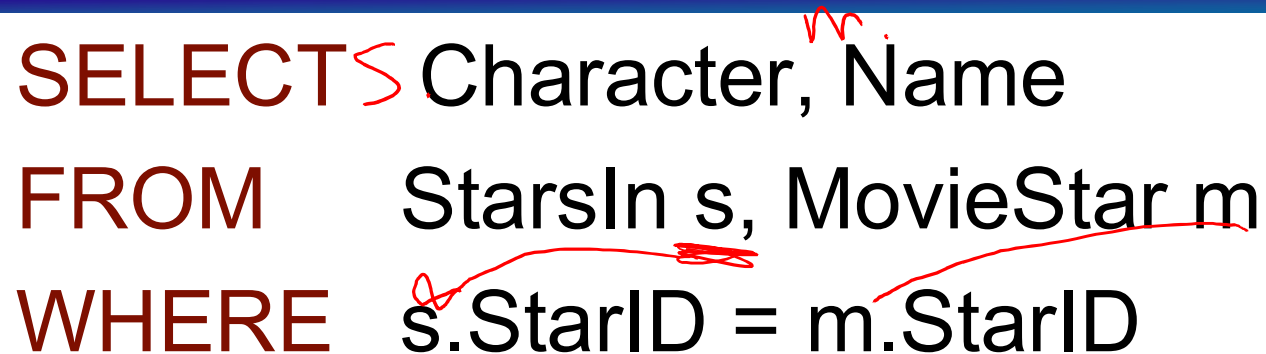
SID	BID	Day
22	101	1998-10-10
B	102	1998-10-10
22	103	1998-10-08
22	104	1998-07-10

Basic SQL/RA comparison example 3

- Find the person names and character names of those who have been in movies
- In order to do this we need to use joins.
How can we do joins in SQL?
 - $\pi \rightarrow$ SELECT clause
 - $\sigma \rightarrow$ WHERE clause
 - $\bowtie \rightarrow$ FROM and WHERE clause

Joins in SQL

```
SELECT Character, Name
FROM StarsIn s, MovieStar m
WHERE s.StarID = m.StarID
```



- Cross product specified by From clause
- Can alias relations (e.g., “StarsIn s”)
- Conditions specified in where clause

Clicker Question: Joins

Consider R :

a	b
0	0
0	1
1	0
1	1

S:

a	b
0	0
0	1
1	0
1	1

T:

a	b
0	0
0	1
1	0
1	1

SELECT R.a, R.b, S.b, T.b
FROM R, S, T
WHERE R.b = S.a AND S.b <> T.b (note: <> == 'not equals')

Compute the results

Which of the following are true:

- A. (0,1,1,0) appears twice.
- B. (1,1,0,1) does not appear.
- C. (1,1,1,0) appears once.
- D. All are true
- E. None are true

Clicker Question: Joins

Consider R :

a	b
0	0
0	1
1	0
1	1

S:

a	b
0	0
0	1
1	0
1	1

T:

a	b
0	0
0	1
1	0
1	1

SELECT R.a, R.b, S.b, T.b
FROM R, S, T
WHERE R.b = S.a AND S.b <> T.b (note: <> == 'not equals')

Compute the results

Which of the following are true:

- A. (0,1,1,0) appears twice.
- B. (1,1,0,1) does not appear.
- C. (1,1,1,0) appears once.
- D. All are true
- E. None are true

True R(0,1) S(1,1), T(0,0)&

R(0,1), S(1,1), T(1,0),

False: R(1,1), S(1,0), T(0,1)

False: like A but use R(1, 1)